Deep Learning & Applied AI

Introduction

Emanuele Rodolà
rodola@di.uniroma1.it

SAPIENZA
UNIVERSITÀ DI ROMA

2nd semester a.y. 2021/2022 · February 23, 2022

Token

574752

Please confirm your presence in the physical classroom using the token above.

Logistics

- **Lecturer:** Prof. Emanuele Rodolà
- **Assistants:** Dr. Luca Moschella and Dr. Donato Crisostomi
  Coding sessions, project support, technicalities

Logistics

- **Lecturer:** Prof. Emanuele Rodolà
- **Assistants:** Dr. Luca Moschella and Dr. Donato Crisostomi
  Coding sessions, project support, technicalities
- **When:** Wednesdays 16:00–**18:30** and Thursdays 10:00–**11:30**

Logistics

- **Lecturer:** Prof. Emanuele Rodolà
- **Assistants:** Dr. Luca Moschella and Dr. Donato Crisostomi
  Coding sessions, project support, technicalities
- **When:** Wednesdays 16:00–**18:30** and Thursdays 10:00–**11:30**
- **Where:**
  Physically: Aula 1 & 2 - Aule L Via del Castro Laurenziano 7a
  Virtually: Zoom, Meeting ID: 475 234 9941, Passcode: 3K7xrM

Logistics

- **Lecturer:** Prof. Emanuele Rodolà
- **Assistants:** Dr. Luca Moschella and Dr. Donato Crisostomi
  Coding sessions, project support, technicalities
- **When:** Wednesdays 16:00–**18:30** and Thursdays 10:00–**11:30**
- **Where:**
  Physically: Aula 1 & 2 - Aule L Via del Castro Laurenziano 7a
  Virtually: Zoom, Meeting ID: 475 234 9941, Passcode: 3K7xrM
- **Office Hours:** Open a discussion on github! (see next slide)
- **Official website:** https://erodola.github.io/DLAI-s2-2022/
  Check frequently for news and material (code, papers, ...)!

Repository & Issues

The course is hosted on Github at the url:

https://github.com/erodola/DLAI-s2-2022

Repository & Issues

The course is hosted on Github at the url:

https://github.com/erodola/DLAI-s2-2022

You can use github to ask questions, in particular:

- Start a discussion to create a new topic / question
  (this replaces the issue system used last year)
- Reply to discussions started by others
- Please use discussions instead of direct emails to the Professor,
  unless you have private reasons.

Repository & Issues

The course is hosted on Github at the url:

https://github.com/erodola/DLAI-s2-2022

You can use github to ask questions, in particular:

- Start a discussion to create a new topic / question
  (this replaces the issue system used last year)
- Reply to discussions started by others
- Please use discussions instead of direct emails to the Professor,
  unless you have private reasons.

Webpage and repository are your main source of information, and replace completely the need for a mailing list. Check them often!

Disclaimer

We are at the third edition of a popular course.

- We will mostly follow the same material as the previous edition (with some semplifications and optional parts)
- The exam will be the same

Disclaimer

We are at the third edition of a popular course.

- We will mostly follow the same material as the previous edition (with some semplifications and optional parts)
- The exam will be the same
- We will alternate between theoretical and lab classes

Disclaimer

We are at the third edition of a popular course.

- We will mostly follow the same material as the previous edition (with some semplifications and optional parts)
- The exam will be the same
- We will alternate between theoretical and lab classes
- We will have guest lecturers on specific topics

Disclaimer

We are at the third edition of a popular course.

- We will mostly follow the same material as the previous edition (with some semplifications and optional parts)
- The exam will be the same
- We will alternate between theoretical and lab classes
- We will have guest lecturers on specific topics
- The course is challenging

Being proficient with Keras, or getting a good grade at other ML courses, will not guarantee success.

24/7 playground

Please join the discord server of the course:

https://discord.gg/emQ9UPeVwA

This will be our meeting place for the lab sessions.

You can use it at any time for the entire semester to organize meetings, collaborate on the lab assignments, etc.

The TAs and myself will be always online during the lab lectures, and will check it daily even on non-lecture days.

Failure cases from past exams

It is not advised to:

Failure cases from past exams

It is not advised to:

- **Trust your memory, intuition, or working knowledge**
  Example: know the general idea of a GAN, but not the formal reasoning behind it.

Failure cases from past exams

It is not advised to:

- **Trust your memory, intuition, or working knowledge**
  Example: know the general idea of a GAN, but not the formal reasoning behind it.
- **Disregard the fundamentals**
  Example: know what convolution means for images, but not being able to write it down as a linear map.

Failure cases from past exams

It is not advised to:

- **Trust your memory, intuition, or working knowledge**
  Example: know the general idea of a GAN, but not the formal reasoning behind it.
- **Disregard the fundamentals**
  Example: know what convolution means for images, but not being able to write it down as a linear map.
- **Ignore the lab material (notebooks)**
  Example: know the theory behind a VAE, but not the reparametrization trick.

Failure cases from past exams

It is not advised to:

- **Trust your memory, intuition, or working knowledge**
  Example: know the general idea of a GAN, but not the formal reasoning behind it.
- **Disregard the fundamentals**
  Example: know what convolution means for images, but not being able to write it down as a linear map.
- **Ignore the lab material (notebooks)**
  Example: know the theory behind a VAE, but not the reparametrization trick.
- **Postpone study**
  New lectures usually assume knowledge from the previous lectures.

Recipe for success

Try to enjoy the course!

Take this as an opportunity to learn in depth.

Ask questions when in doubt.

Who am I?

- Had research positions at U Tokyo, TU Munich, U Lugano and visiting positions at Harvard, Stanford, Ecole polytechnique, Technion among others
- Research: digital geometry processing, geometric deep learning
- Team: ∼25 members from physics, engineering, computer science **GLADIA** group of Geometry, Learning and AI
- If you have ideas, approach us for projects / theses

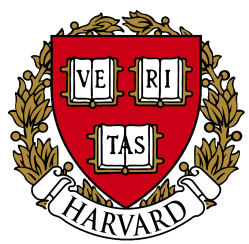


nature.com/nmeth February 2020 Vol. 17 No. 2

nature methods

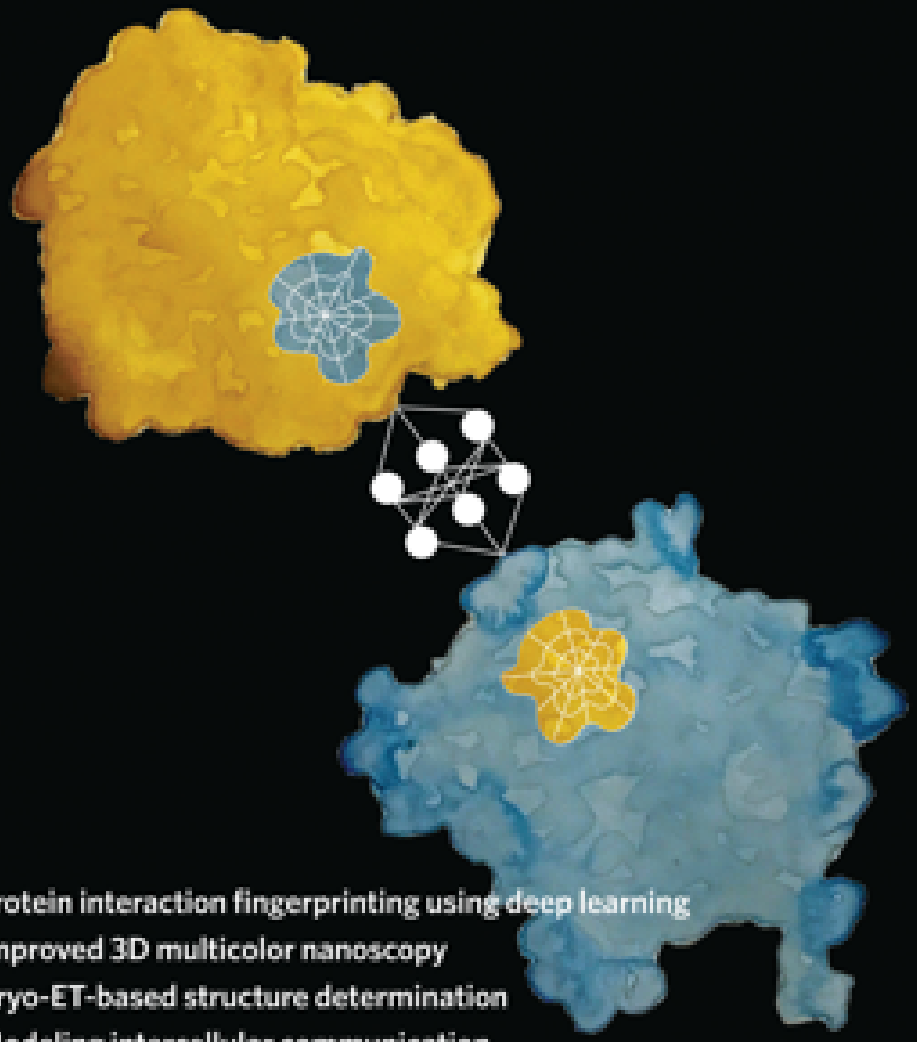

Protein interaction fingerprinting using deep learning

Improved 3D multicolor nanoscopy

Cryo-ET-based structure determination

Modeling intercellular communication

The Bioconductor project for single-cell analysis

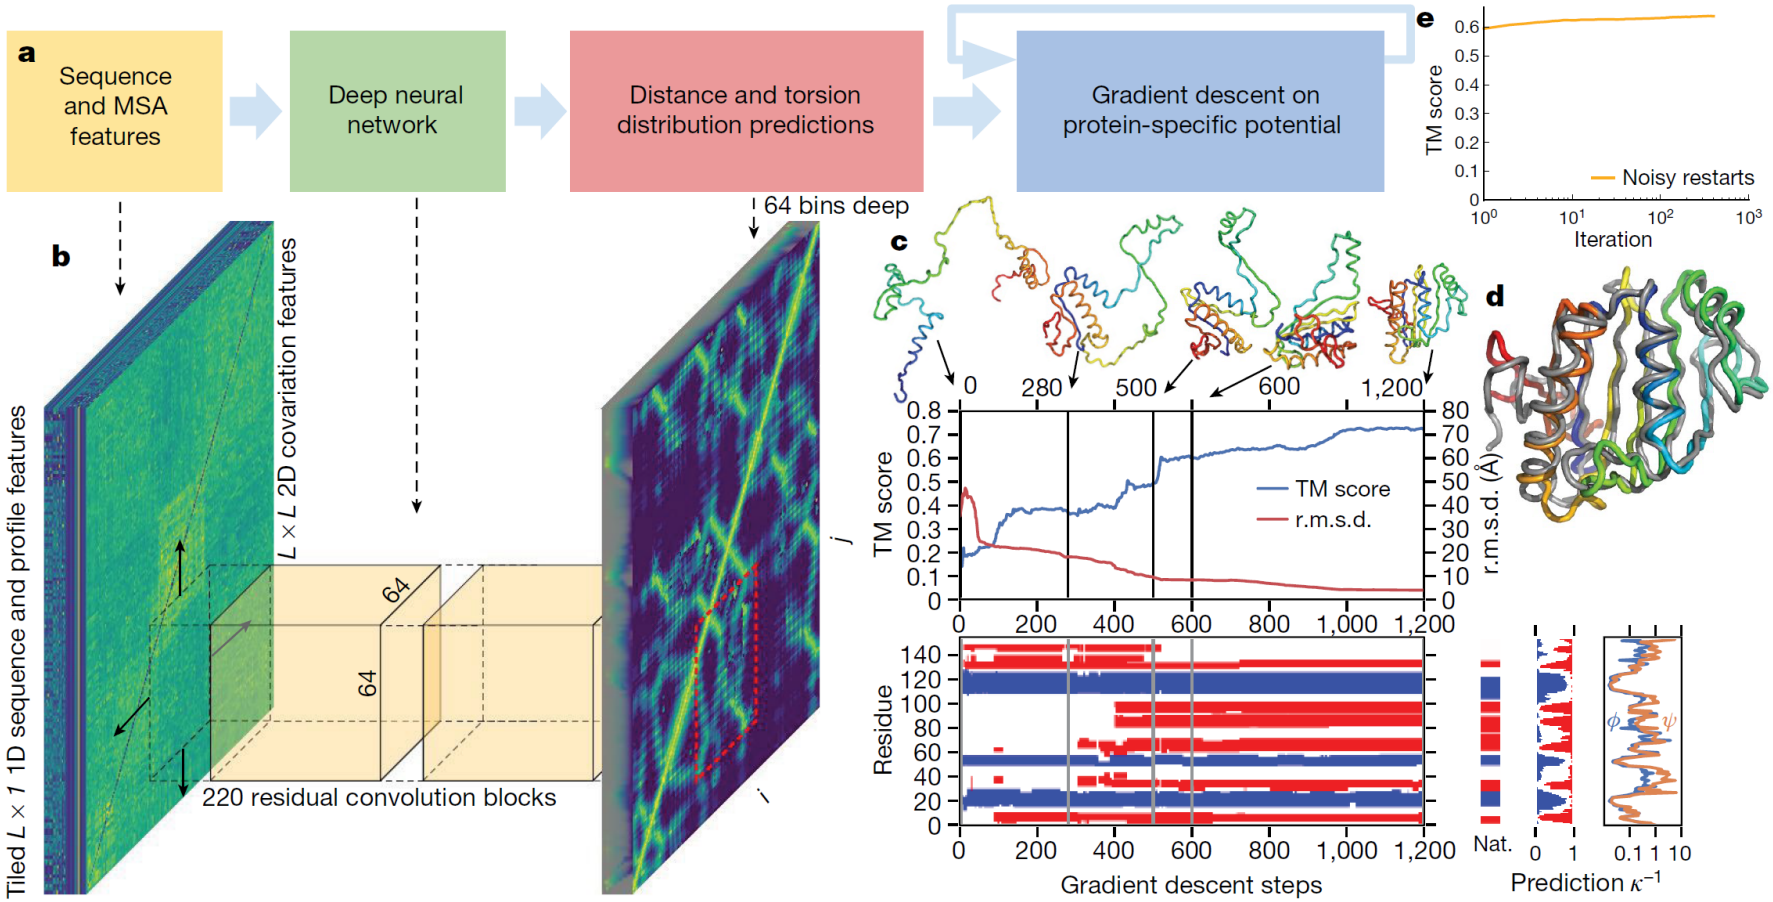

https://deepmind.com/research/case-studies/alphafold

Pre-requisites and reading material

No official textbook.

Specific references will be given throughout the course in the form of book chapters and scientific articles.

Pre-requisites and reading material

No official textbook.

Specific references will be given throughout the course in the form of book chapters and scientific articles.

Warning: A blog post does not (always) count as a reliable reference, since it's not peer-reviewed.

Pre-requisites and reading material

No official textbook.

Specific references will be given throughout the course in the form of book chapters and scientific articles.

Warning: A blog post does not (always) count as a reliable reference, since it's not peer-reviewed.

Pre-requisites:

- **Programming fundamentals**. We will use Python
- Welcome (not mandatory): linear algebra, calculus

Pre-requisites and reading material

Main source: specialized conferences and journals:

NeurIPS, ICML, ICLR, JMLR, CVPR, etc.

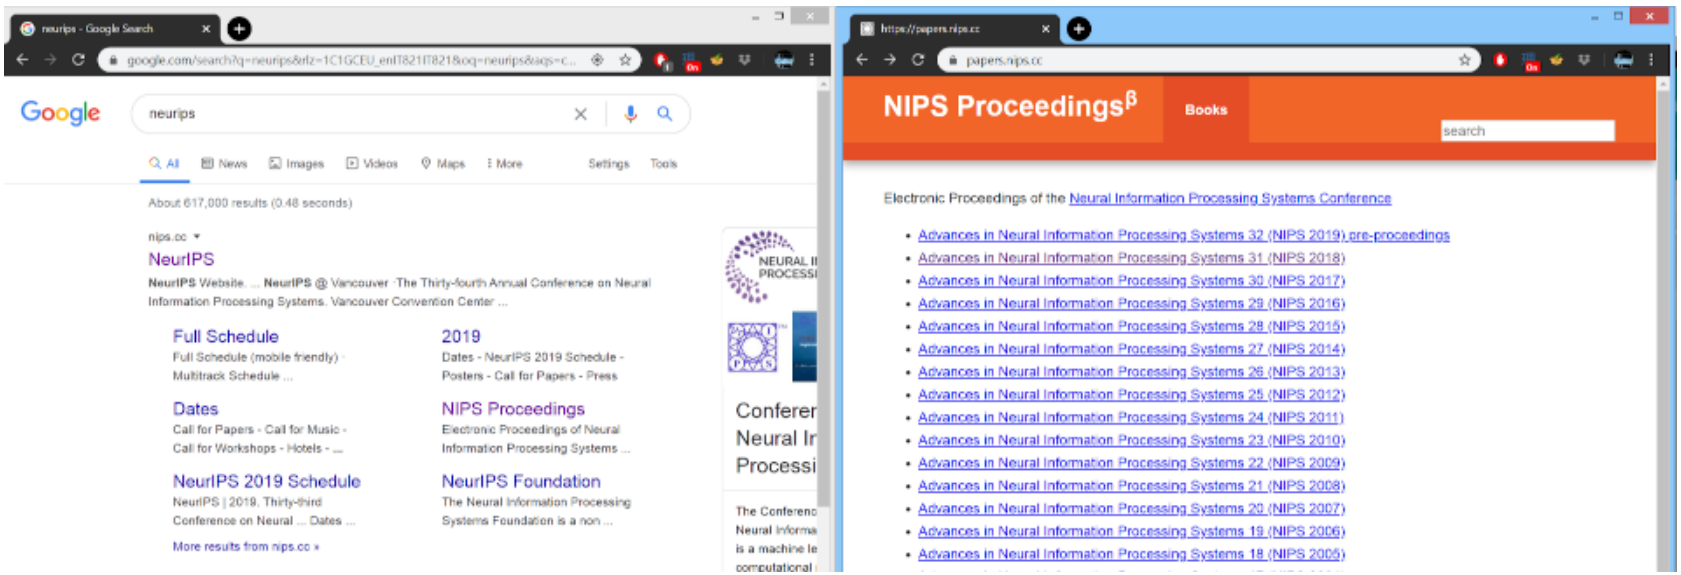

Pre-requisites and reading material

Main source: specialized conferences and journals:

NeurIPS, ICML, ICLR, JMLR, CVPR, etc.

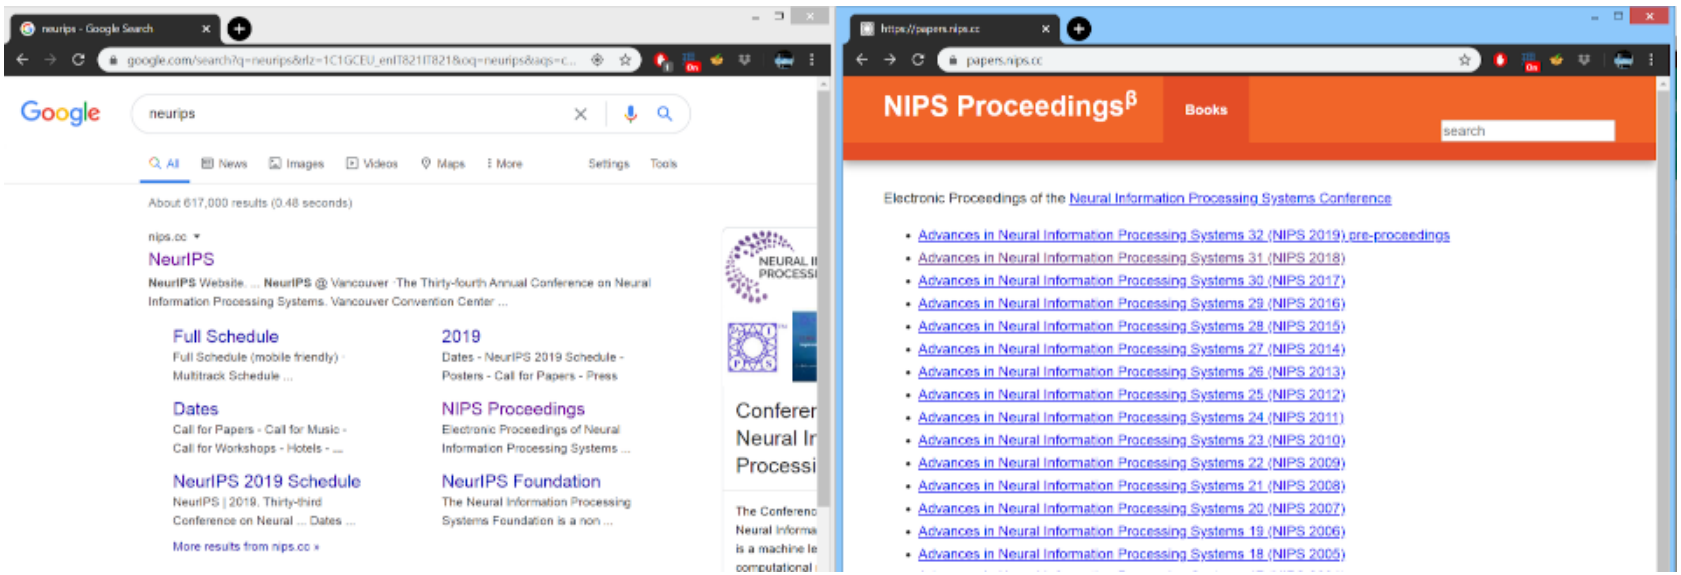

Applications: we'll mostly illustrate things in the areas of computer vision, geometry processing, graphics, social networks, biochemistry.

Grading

1. Midterm self-evaluation (*not graded*)
   When: first half of April

Grading

1. Midterm self-evaluation (*not graded*)
   When: first half of April
2. Project with report (fixed template)
   When: Within the academic year

   The project can be selected from a pool, or can be proposed.

Grading

1. Midterm self-evaluation (*not graded*)
   When: first half of April
2. Project with report (fixed template)
   When: Within the academic year

   The project can be selected from a pool, or can be proposed.
3. Written exam
   When: June 10th and July 8th

   Some past exams are on the course webpage.

Grading

1. Midterm self-evaluation (*not graded*)
   When: first half of April
2. Project with report (fixed template)
   When: Within the academic year

   The project can be selected from a pool, or can be proposed.
3. Written exam
   When: June 10th and July 8th

   Some past exams are on the course webpage.
4. Final grade is 40% project and 60% written exam

Grading

1. Midterm self-evaluation (*not graded*)
   When: first half of April
2. Project with report (fixed template)
   When: Within the academic year

   The project can be selected from a pool, or can be proposed.
3. Written exam
   When: June 10th and July 8th

   Some past exams are on the course webpage.
4. Final grade is 40% project and 60% written exam
5. Optional: oral exam, contributing ± 3 points to the final grade

Grading

1. Midterm self-evaluation (*not graded*)
   When: first half of April
2. Project with report (fixed template)
   When: Within the academic year

   The project can be selected from a pool, or can be proposed.
3. Written exam
   When: June 10th and July 8th

   Some past exams are on the course webpage.
4. Final grade is 40% project and 60% written exam
5. Optional: oral exam, contributing ± 3 points to the final grade

In class, be prepared:

- Download/print the slides beforehand
- Take notes: not everything will be on the slides
- Bring your laptop: we'll do live coding sessions

What will you get out of this course?
(if you study)

- You will acquire solid fundamental skills for understanding, analyzing, and designing deep learning models

What will you get out of this course?
(if you study)

- You will acquire solid fundamental skills for understanding, analyzing, and designing deep learning models
- You will be able to grasp and elaborate on more advanced topics published in the top machine learning venues

What will you get out of this course?
(if you study)

- You will acquire solid fundamental skills for understanding, analyzing, and designing deep learning models
- You will be able to grasp and elaborate on more advanced topics published in the top machine learning venues
- You will get practical development expertise on applied problems


Mathematical tools

- Linear algebra
- Calculus
- Optimization
- Discrete mathematics
- Probability & statistics
- Metric and differential geometry

This is not an easy course, but results will speak for themselves.

We will have to develop ways to evaluate, visualize, and quantify what we are doing. Going blind-folded and regarding learning models as black boxes will not bring us very far!

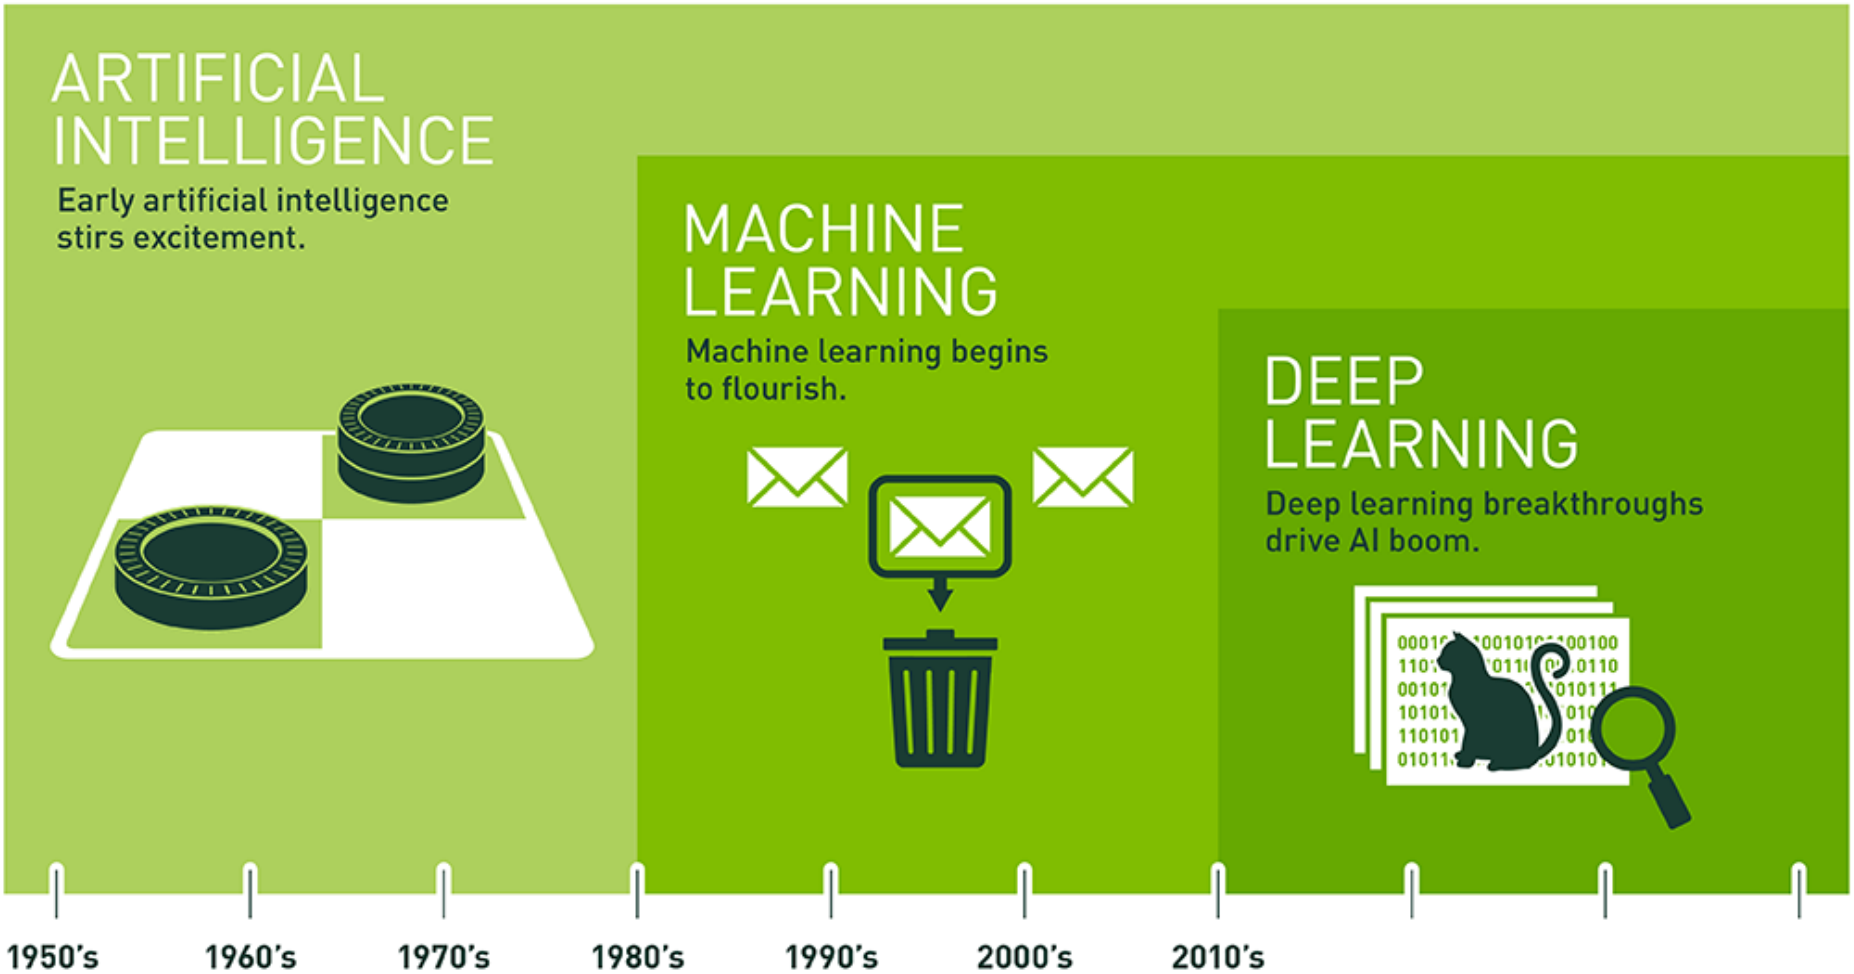

Since an early flush of optimism in the 1950s, smaller subsets of artificial intelligence – first machine learning, then deep learning, a subset of machine learning – have created ever larger disruptions.

Image: Michael Copeland, NVIDIA

Deep learning is everywhere

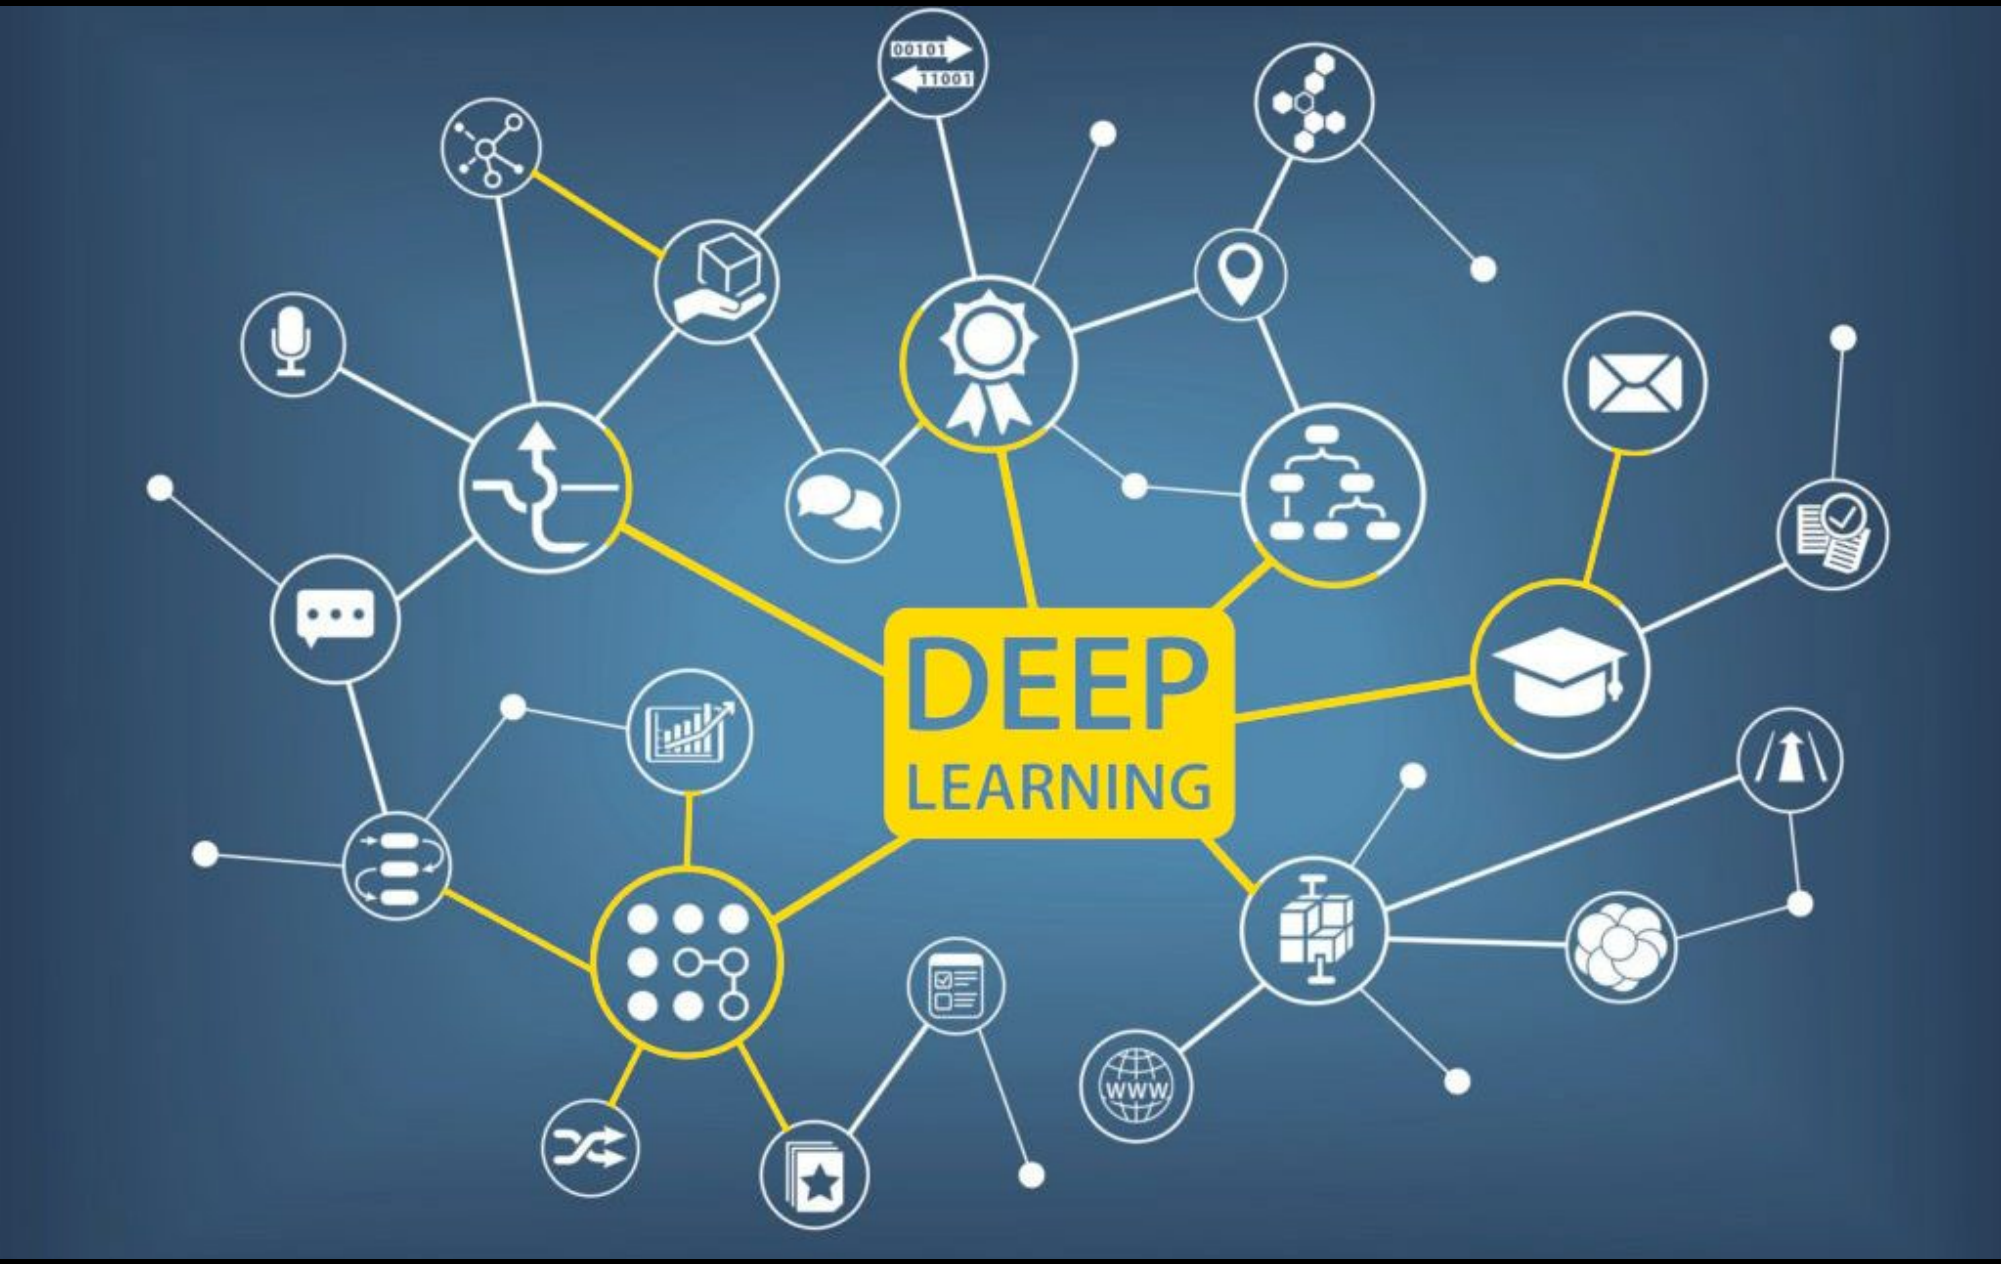

Showcase

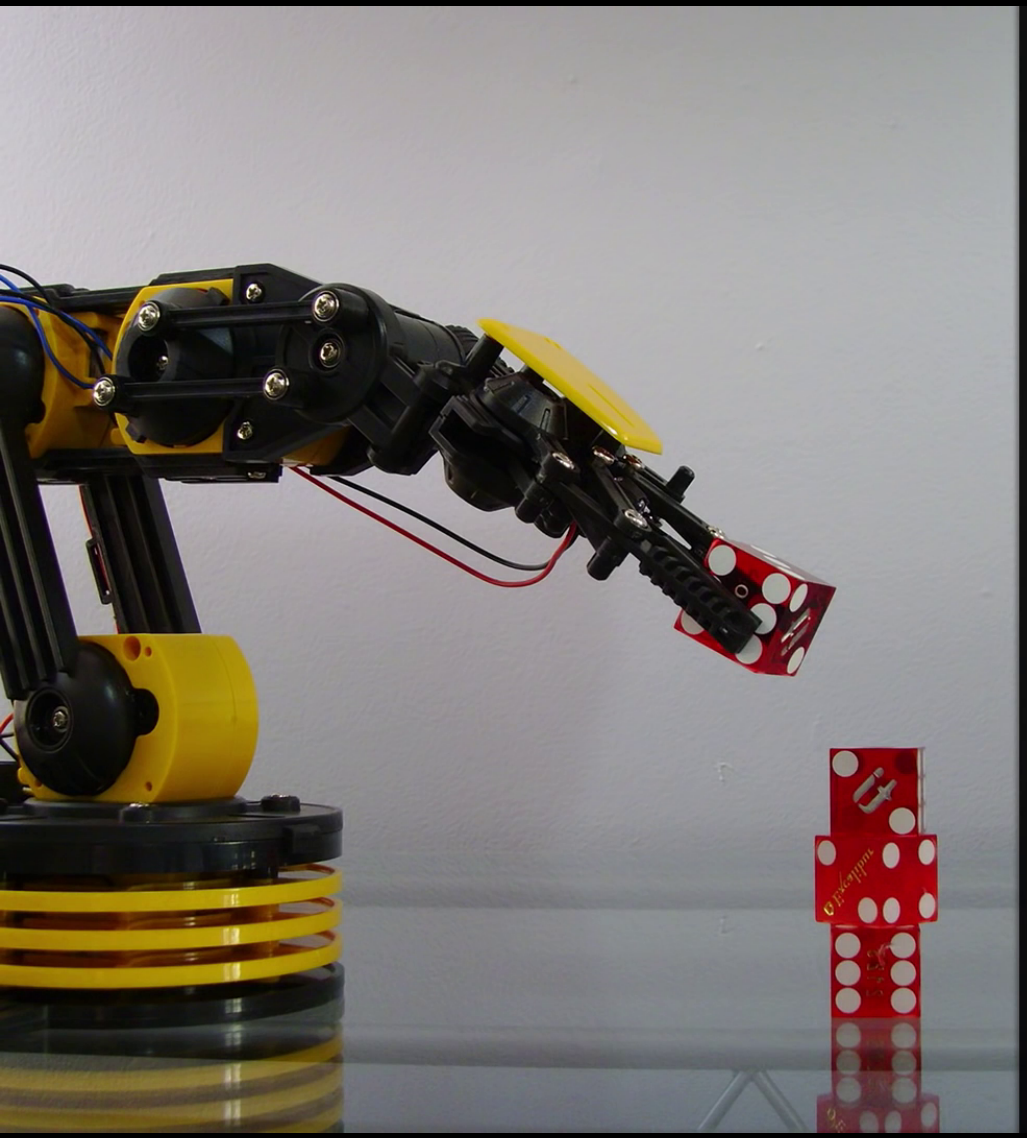

TWO MINUTE PAPERS

WITH KÁROLY ZSOLNAI-FEHÉR (KZF)

10 MORE COOL DEEP LEARNING APPLICATIONS

Disclaimer: I was not part of this research project,
I am merely providing commentary on this.