

Deep Learning & Applied AI

Regularization, batchnorm and dropout

Emanuele Rodolà
rodola@di.uniroma1.it



SAPIENZA
UNIVERSITÀ DI ROMA

Regularization

Overfitting often happens with limited training data:

parameters \gg # training examples

Regularization

Overfitting often happens with limited training data:

$$\# \text{ parameters} \gg \# \text{ training examples}$$

Regularization is a general mechanism to reduce overfitting and thus improve generalization.

Regularization

Overfitting often happens with limited training data:

$$\# \text{ parameters} \gg \# \text{ training examples}$$

Regularization is a general mechanism to reduce overfitting and thus improve generalization.

General idea: reduce the number of free parameters.

Regularization

Overfitting often happens with limited training data:

$$\# \text{ parameters} \gg \# \text{ training examples}$$

Regularization is a general mechanism to reduce overfitting and thus improve generalization.

General idea: reduce the number of free parameters.

- Eliminate network weights.

e.g. estimate network sensitivity w.r.t. each weight.

Regularization

Overfitting often happens with limited training data:

$$\# \text{ parameters} \gg \# \text{ training examples}$$

Regularization is a general mechanism to reduce overfitting and thus improve generalization.

General idea: reduce the number of free parameters.

- Eliminate network weights.
e.g. estimate network sensitivity w.r.t. each weight.
- Weight sharing (i.e. # weights < # connections).

Regularization

Overfitting often happens with limited training data:

$$\# \text{ parameters} \gg \# \text{ training examples}$$

Regularization is a general mechanism to reduce overfitting and thus improve generalization.

General idea: reduce the number of free parameters.

- Eliminate network weights.
e.g. estimate network sensitivity w.r.t. each weight.
- Weight sharing (i.e. # weights < # connections).
- Explicit penalties.

Regularization

Overfitting often happens with limited training data:

$$\# \text{ parameters} \gg \# \text{ training examples}$$

Regularization is a general mechanism to reduce overfitting and thus improve generalization.

General idea: reduce the number of free parameters.

- Eliminate network weights.
e.g. estimate network sensitivity w.r.t. each weight.
- Weight sharing (i.e. # weights < # connections).
- Explicit penalties.
- Implicit regularization.

Regularization

Any modification that is intended to reduce the generalization error but not the training error.

Weight penalties

$$\underbrace{\ell(\Theta)}_{\text{loss}} + \lambda \underbrace{\rho(\Theta)}_{\text{regularizer}}$$

Weight penalties

$$\underbrace{\ell(\Theta)}_{\text{loss}} + \lambda \underbrace{\rho(\Theta)}_{\text{regularizer}}$$

The regularizer induces a trade-off:

data **fidelity** vs. model **complexity**

Weight penalties

$$\underbrace{\ell(\Theta)}_{\text{loss}} + \lambda \underbrace{\rho(\Theta)}_{\text{regularizer}}$$

The regularizer induces a trade-off:

data **fidelity** vs. model **complexity**

Typical penalties:

- Tikhonov (L_2) regularization \Rightarrow promotes **shrinkage**

Weight penalties

$$\underbrace{\ell(\Theta)}_{\text{loss}} + \lambda \underbrace{\rho(\Theta)}_{\text{regularizer}}$$

The regularizer induces a trade-off:

data **fidelity** vs. model **complexity**

Typical penalties:

- Tikhonov (L_2) regularization \Rightarrow promotes **shrinkage**
- Lasso (L_1) regularization \Rightarrow promotes **sparsity** or **weight selection**

Weight penalties

$$\underbrace{\ell(\Theta)}_{\text{loss}} + \lambda \underbrace{\rho(\Theta)}_{\text{regularizer}}$$

The regularizer induces a trade-off:

data **fidelity** vs. model **complexity**

Typical penalties:

- Tikhonov (L_2) regularization \Rightarrow promotes **shrinkage**
- Lasso (L_1) regularization \Rightarrow promotes **sparsity** or **weight selection**
- Bounded L_2 norm at each layer $\|\mathbf{W}^{(\ell)}\|_F \leq c^{(\ell)}$

Weight penalties

$$\underbrace{\ell(\Theta)}_{\text{loss}} + \lambda \underbrace{\rho(\Theta)}_{\text{regularizer}}$$

The regularizer induces a trade-off:

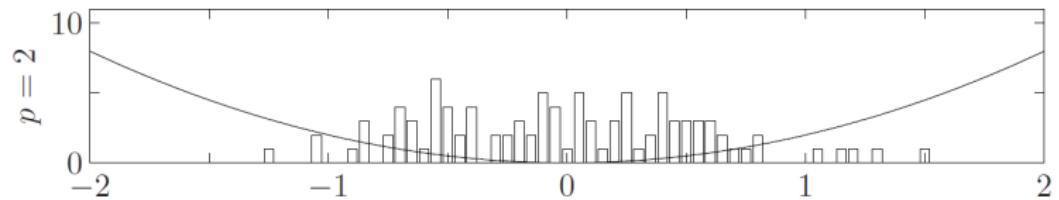
data **fidelity** vs. model **complexity**

Typical penalties:

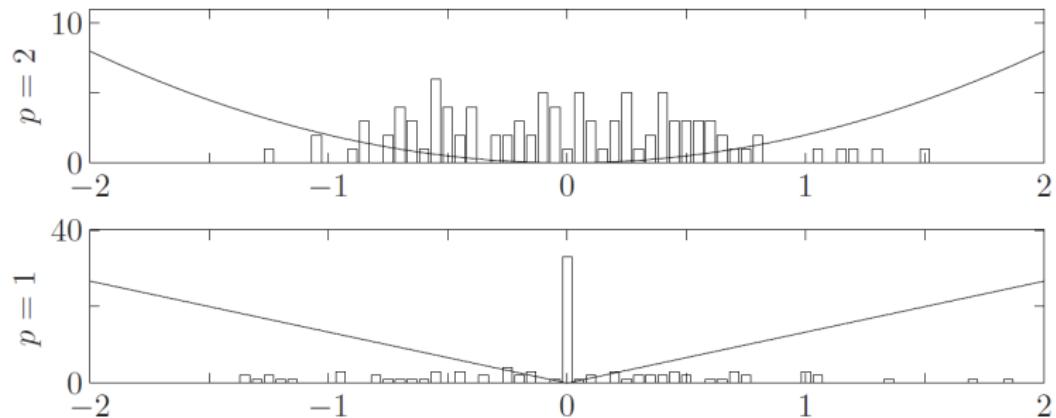
- Tikhonov (L_2) regularization \Rightarrow promotes **shrinkage**
- Lasso (L_1) regularization \Rightarrow promotes **sparsity** or **weight selection**
- Bounded L_2 norm at each layer $\|\mathbf{W}^{(\ell)}\|_F \leq c^{(\ell)}$

After training, the L_p magnitude of each weight reflects its importance.

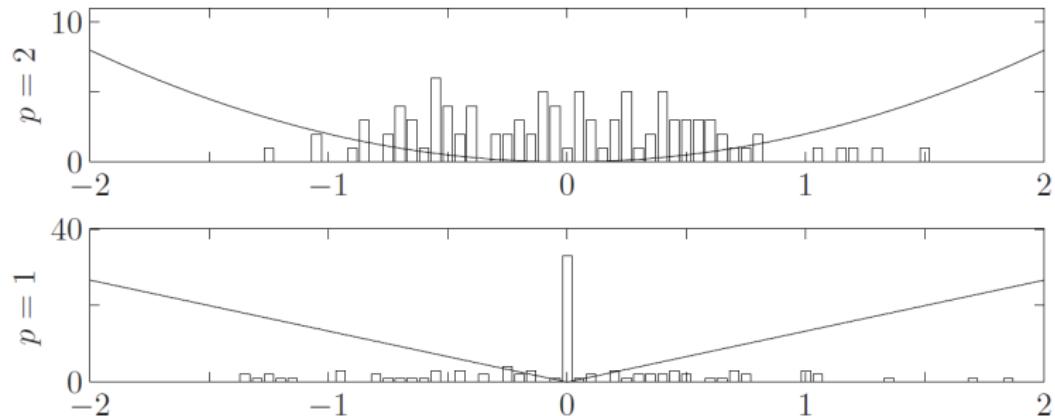
L_2 vs L_1 penalties



L_2 vs L_1 penalties

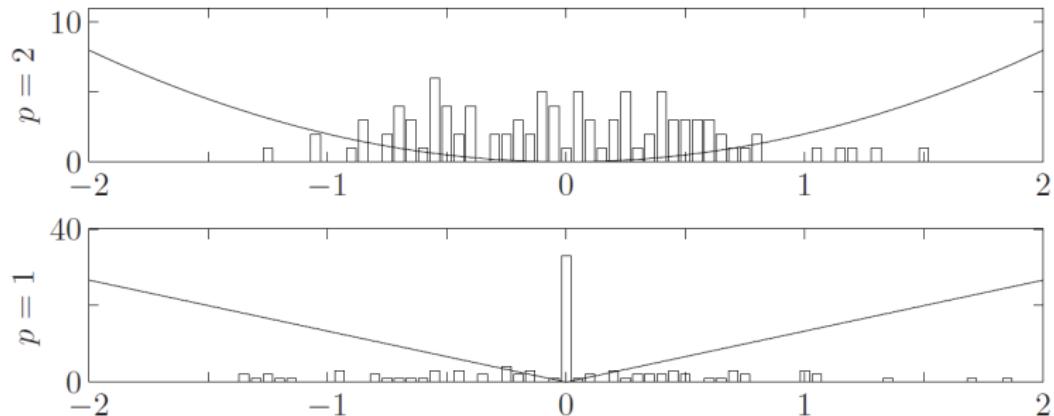


L_2 vs L_1 penalties



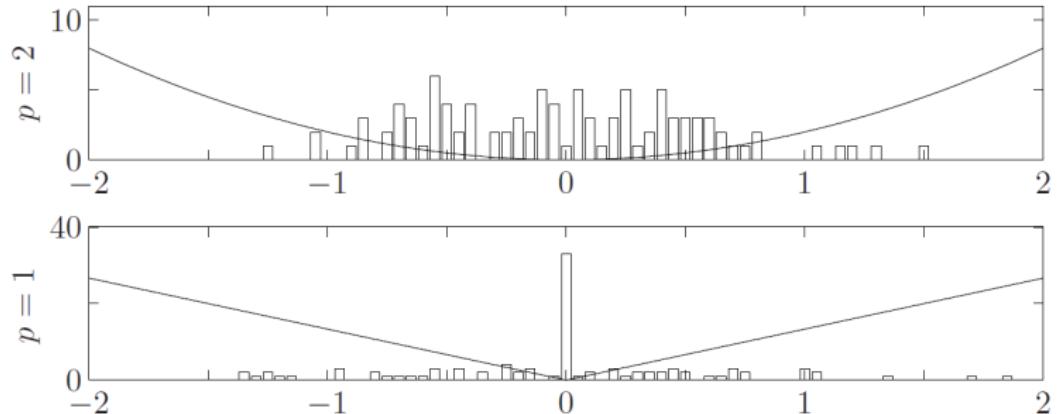
- Big reduction in $\|\Theta\|_2$ if you scale down the values > 1

L_2 vs L_1 penalties



- Big reduction in $\|\Theta\|_2$ if you scale down the values > 1
- Almost no reduction in $\|\Theta\|_2$ for values < 1 . Sparsity is discouraged!

L_2 vs L_1 penalties

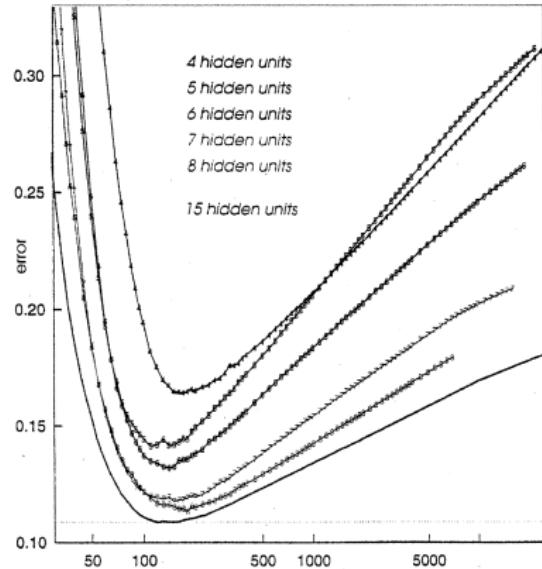


- Big reduction in $\|\Theta\|_2$ if you scale down the values > 1
- Almost no reduction in $\|\Theta\|_2$ for values < 1 . Sparsity is discouraged!
- All the values are treated the same in $\|\Theta\|_1$, no matter if they are > 1 or < 1 . Any value can be set to zero, leading to **sparse solutions**.

Source code: <https://github.com/ievron/RegularizationAnimation/>

Detecting overfitting

Overfitting can be recognized by looking at the **validation error**:

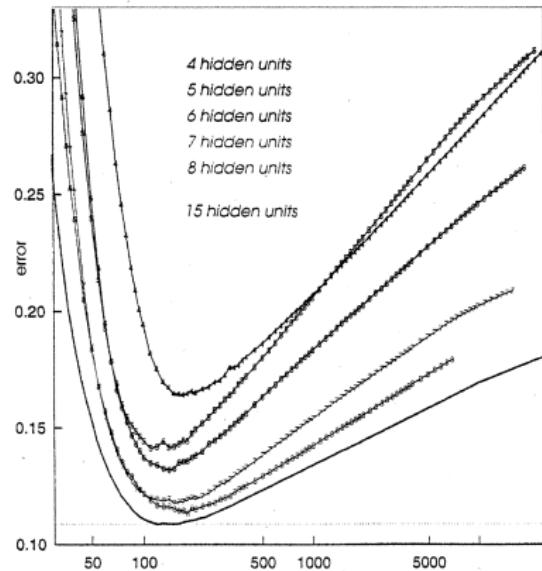


Weigend, "On overfitting and the effective number of hidden units", 1993

Detecting overfitting

Overfitting can be recognized by looking at the validation error:

- Small networks can also overfit.

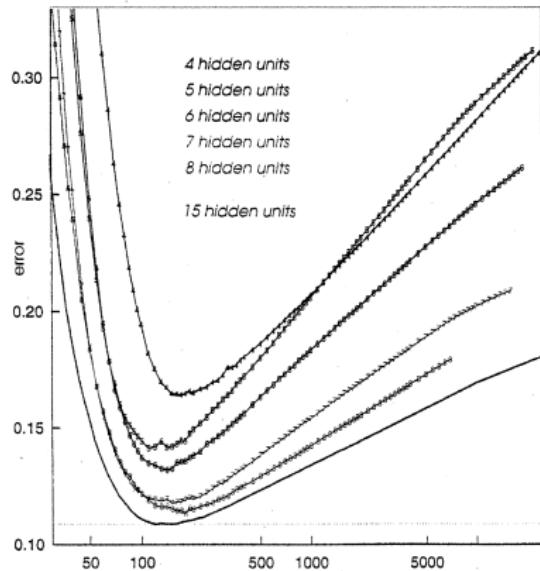


Weigend, "On overfitting and the effective number of hidden units", 1993

Detecting overfitting

Overfitting can be recognized by looking at the validation error:

- Small networks can also overfit.
- Large networks have best performance if they stop early.

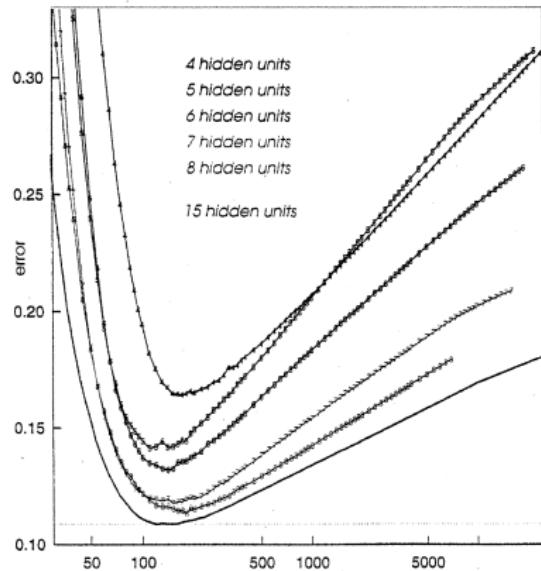


Weigend, "On overfitting and the effective number of hidden units", 1993

Detecting overfitting: Early stopping

Overfitting can be recognized by looking at the validation error:

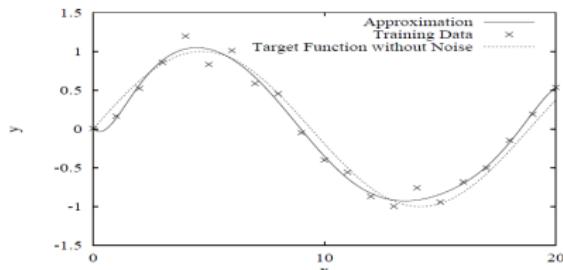
- Small networks can also overfit.
- Large networks have best performance if they stop early.
- Early stopping: Stop training as soon as performance on a validation set decreases.



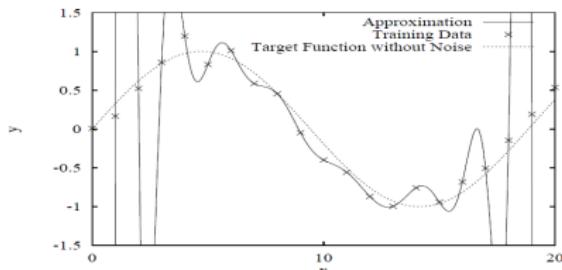
Weigend, "On overfitting and the effective number of hidden units", 1993

Many parameters \neq overfitting

Typical overfitting with polynomial regression:



Order 10

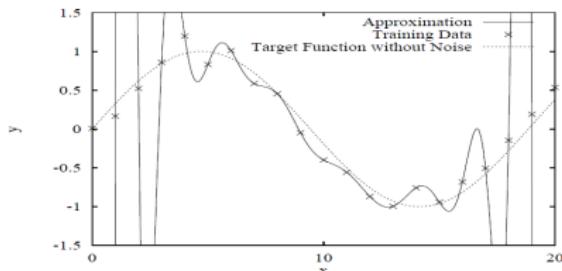
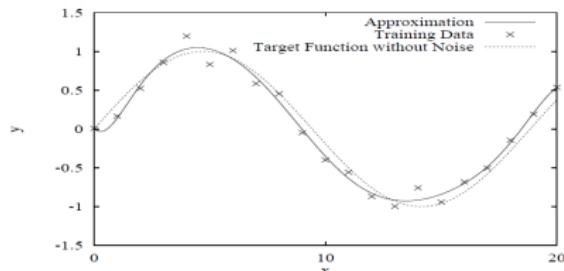


Order 20

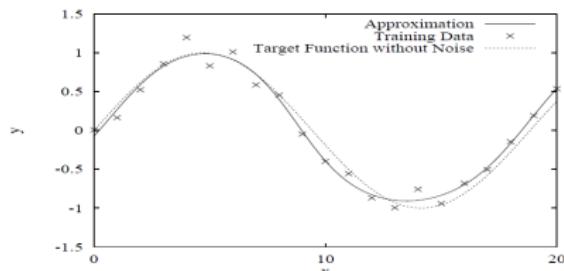
Caruana et al, "Overfitting in Neural Nets: Backpropagation, Conjugate Gradient, and Early Stopping", NIPS 2001

Many parameters \neq overfitting

...but more MLP parameters not always lead to overfitting:

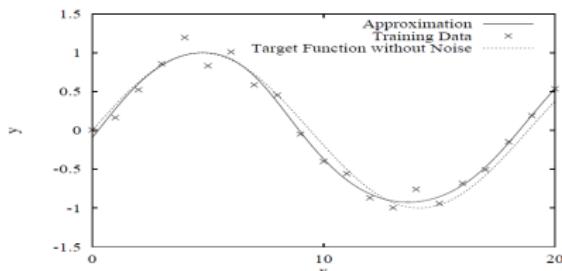


Order 10



10 Hidden Nodes

Order 20

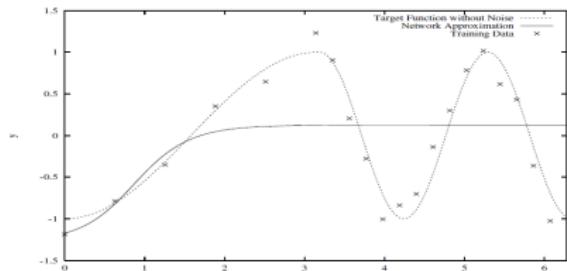


50 Hidden Nodes

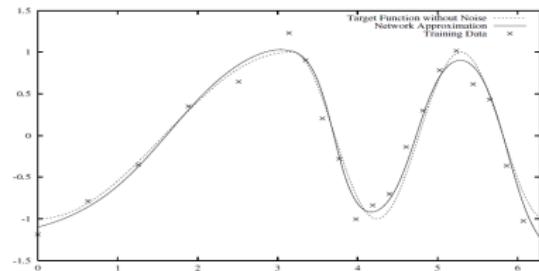
Caruana et al, "Overfitting in Neural Nets: Backpropagation, Conjugate Gradient, and Early Stopping", NIPS 2001

Many parameters \neq overfitting

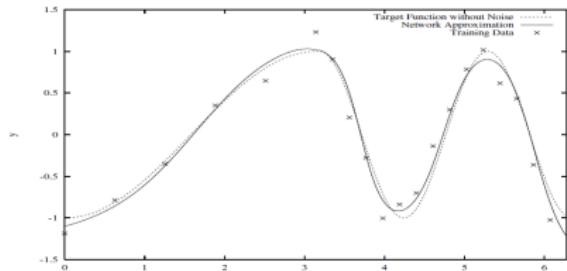
Good fit over all the different data regions:



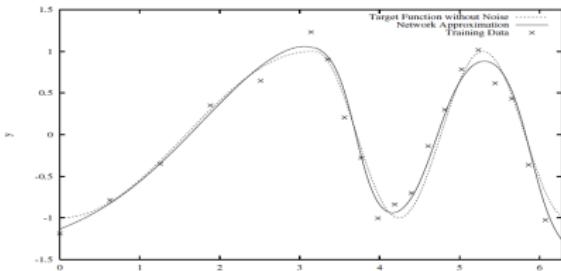
1 Hidden Unit



4 Hidden Units



10 Hidden Units

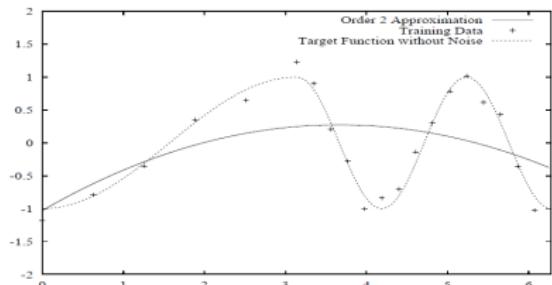


100 Hidden Units

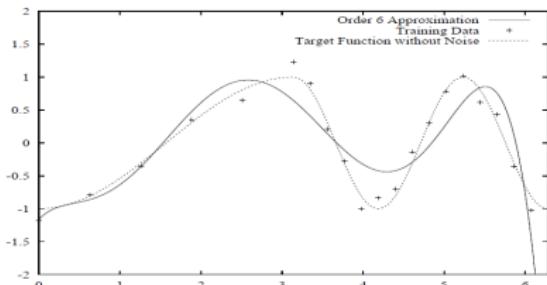
Caruana et al, "Overfitting in Neural Nets: Backpropagation, Conjugate Gradient, and Early Stopping", NIPS 2001

Overfitting as a local phenomenon

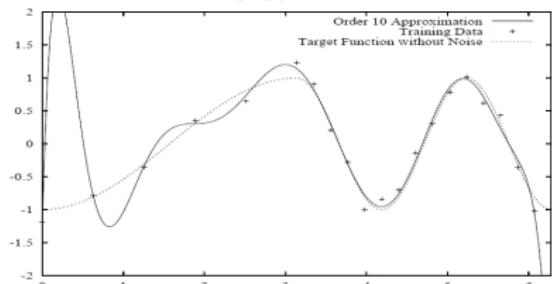
Overfitting is **local** and can vary significantly in different regions:



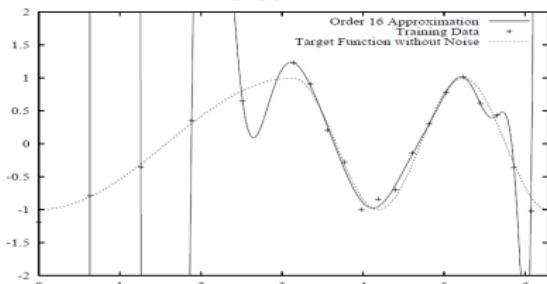
Order 2



Order 6



Order 10

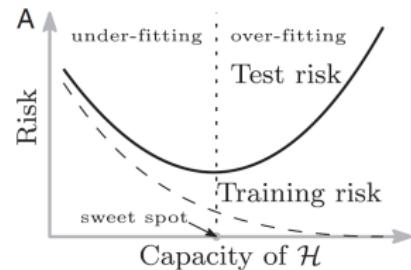


Order 16

Caruana et al, "Overfitting in Neural Nets: Backpropagation, Conjugate Gradient, and Early Stopping", NIPS 2001

Double descent

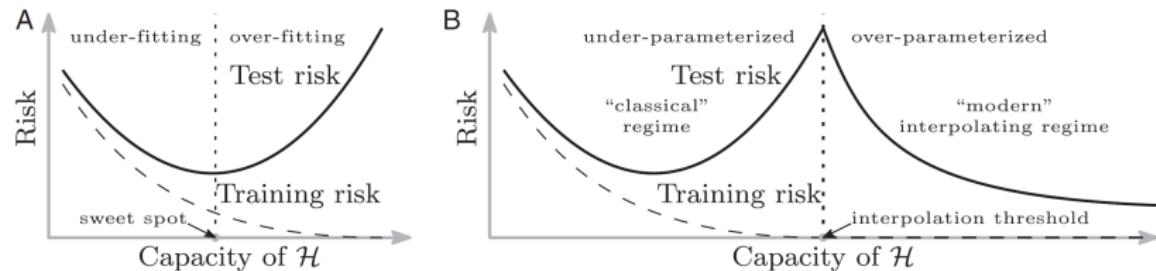
U-shaped curve as a function of # network parameters:



Belkin et al, "Reconciling modern machine-learning practice and the classical bias-variance trade-off", PNAS 2019

Double descent

U-shaped curve as a function of # network parameters:

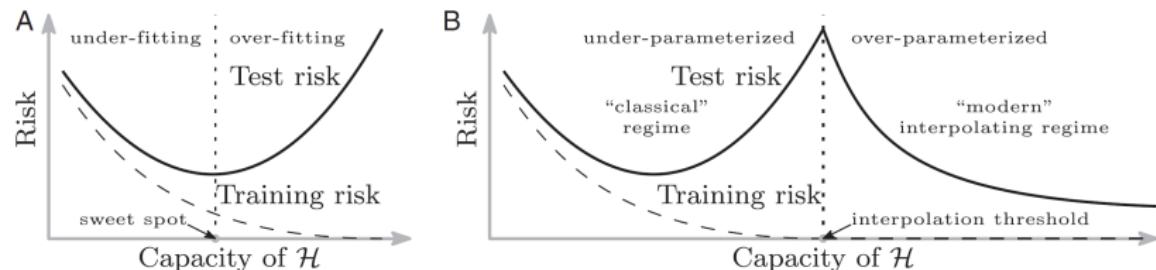


Interpolation: perfect fit on the training data.

Belkin et al, "Reconciling modern machine-learning practice and the classical bias-variance trade-off", PNAS 2019

Double descent

U-shaped curve as a function of # network parameters:



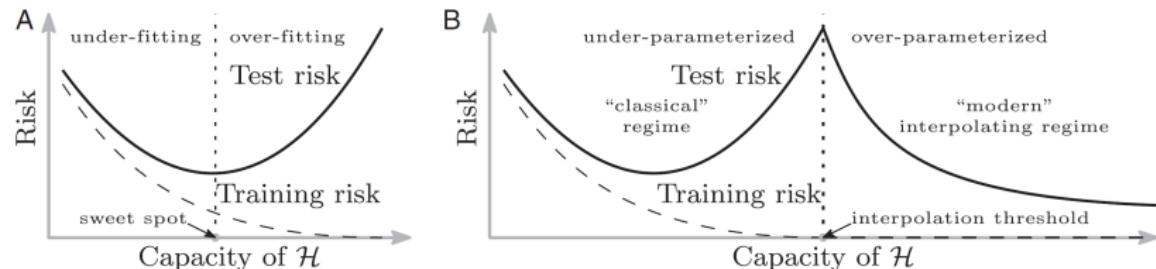
Interpolation: perfect fit on the training data.

“By considering larger function classes, which contain more candidate predictors compatible with the data, we are able to find interpolating functions that have smaller norm and are thus “simpler.” Thus, increasing function class capacity improves performance of classifiers.”

Belkin et al, “Reconciling modern machine-learning practice and the classical bias-variance trade-off”, PNAS 2019

Double descent

U-shaped curve as a function of # network parameters:



Interpolation: perfect fit on the training data.

“By considering larger function classes, which contain more candidate predictors compatible with the data, we are able to find interpolating functions that have smaller norm and are thus “simpler.” Thus, increasing function class capacity improves performance of classifiers.”

The surprising fact is that SGD is able to find such good models.

Belkin et al, “Reconciling modern machine-learning practice and the classical bias-variance trade-off”, PNAS 2019

Early stopping

Early stopping is based on the “smoothness” heuristic:

Representational power **grows** with training time

Caruana et al, “Overfitting in Neural Nets: Backpropagation, Conjugate Gradient, and Early Stopping”, NIPS 2001

Early stopping

Early stopping is based on the “smoothness” heuristic:

Representational power grows with training time

- Initialize with small weights.

Caruana et al, “Overfitting in Neural Nets: Backpropagation, Conjugate Gradient, and Early Stopping”, NIPS 2001

Early stopping

Early stopping is based on the “smoothness” heuristic:

Representational power grows with training time

- Initialize with small weights.
- Simple hypotheses are considered before complex hypotheses.

Caruana et al, “Overfitting in Neural Nets: Backpropagation, Conjugate Gradient, and Early Stopping”, NIPS 2001

Early stopping

Early stopping is based on the “smoothness” heuristic:

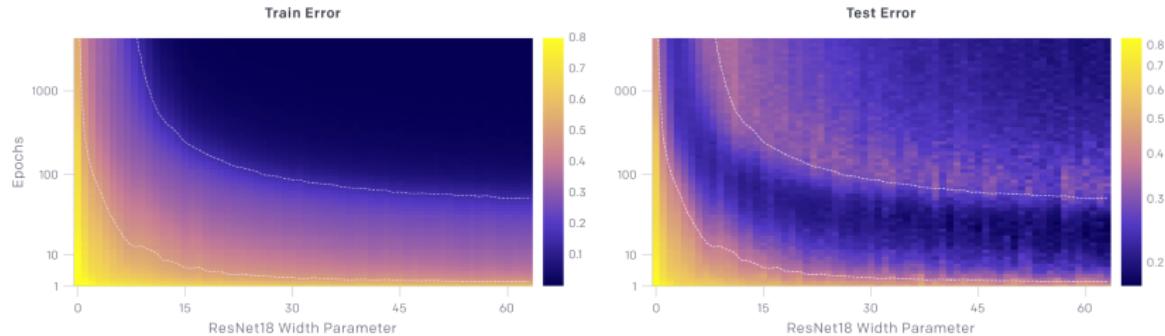
Representational power grows with training time

- Initialize with small weights.
- Simple hypotheses are considered before complex hypotheses.
- Training first explores models similar to what a **smaller net of optimal size** would have learned.

Caruana et al, “Overfitting in Neural Nets: Backpropagation, Conjugate Gradient, and Early Stopping”, NIPS 2001

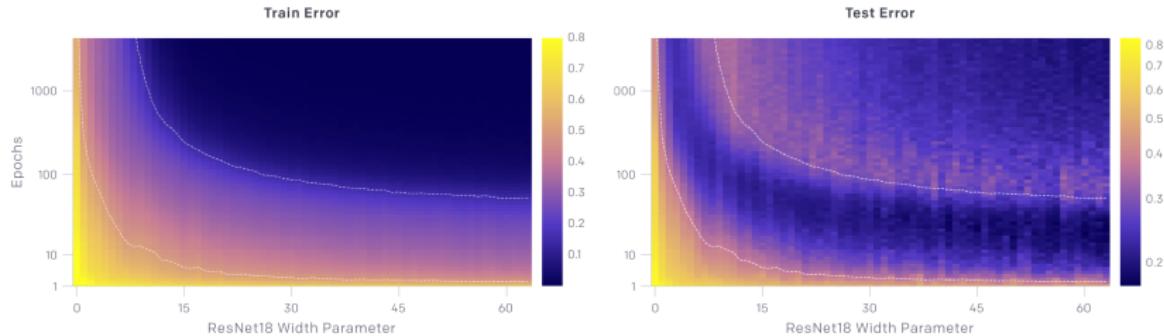
Epoch-wise double descent

There is a regime where training longer reverses overfitting.



Epoch-wise double descent

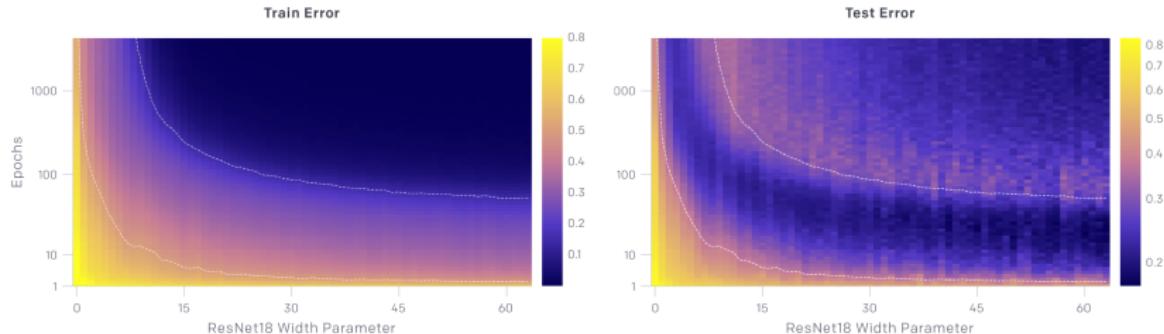
There is a regime where training longer reverses overfitting.



For a fixed number of epochs, the “usual” double descent.

Epoch-wise double descent

There is a regime where training longer reverses overfitting.



For a fixed number of epochs, the “usual” double descent.

For a fixed number of parameters, we observe double descent **as a function of training time**.

Batch normalization

$$\mathbf{x}^{(k)} = \sigma \left(\mathbf{W}^{(k)} \mathbf{x}^{(k-1)} \right)$$

Batch normalization

$$\mathbf{x}^{(k)} = \sigma \left(\mathbf{W}^{(k)} \mathbf{x}^{(k-1)} \right)$$

Internal covariate shift: The input distribution changes at each layer, and the layers need to continuously adapt to the new distribution.

Shimodaira, "Improving predictive inference under covariate shift by weighting the log-likelihood function", 2000

Batch normalization

$$\mathbf{x}^{(k)} = \sigma \left(\mathbf{W}^{(k)} \mathbf{x}^{(k-1)} \right)$$

Internal covariate shift: The input distribution changes at each layer, and the layers need to continuously adapt to the new distribution.

Normalize the features by the statistics computed within the training set:

$$\hat{\mathbf{x}}^{(k)} = \text{normalize}(\mathbf{x}^{(k)}, \mathcal{X})$$

where both \mathbf{x} and \mathcal{X} depend on \mathbf{W} .

Shimodaira, "Improving predictive inference under covariate shift by weighting the log-likelihood function", 2000; Ioffe and Szegedy, "Batch normalization: Accelerating deep network training by reducing internal covariate shift", ICML 2015

Batch normalization

$$\mathbf{x}^{(k)} = \sigma \left(\mathbf{W}^{(k)} \mathbf{x}^{(k-1)} \right)$$

Internal covariate shift: The input distribution changes at each layer, and the layers need to continuously adapt to the new distribution.

Normalize the features by the statistics computed within the training set:

$$\hat{\mathbf{x}}^{(k)} = \text{normalize}(\mathbf{x}^{(k)}, \mathcal{X})$$

where both \mathbf{x} and \mathcal{X} depend on \mathbf{W} .

In particular, backprop will need the partial derivatives:

$$\frac{\partial}{\partial \mathbf{x}} \text{normalize}(\mathbf{x}, \mathcal{X}), \quad \frac{\partial}{\partial \mathcal{X}} \text{normalize}(\mathbf{x}, \mathcal{X})$$

Shimodaira, "Improving predictive inference under covariate shift by weighting the log-likelihood function", 2000; Ioffe and Szegedy, "Batch normalization: Accelerating deep network training by reducing internal covariate shift", ICML 2015

Batch normalization: Transformation

For each dimension of \mathbf{x} , transform:

$$x_i \mapsto \frac{x_i - \mathbb{E}[x_i]}{\sqrt{\text{var}(x_i)}}$$

where mean and variance are computed over the training set.

After the transformation, we get $\text{mean} = 0$ and $\text{var} = 1$.

Ioffe and Szegedy, "Batch normalization: Accelerating deep network training by reducing internal covariate shift", ICML 2015

Batch normalization: Transformation

For each dimension of \mathbf{x} , transform:

$$x_i \mapsto \frac{x_i - \mathbb{E}[x_i]}{\sqrt{\text{var}(x_i)}}$$

where mean and variance are computed over the training set.

After the transformation, we get $\text{mean} = 0$ and $\text{var} = 1$.

Furthermore, introduce **trainable** weights:

$$x_i \mapsto \gamma_i \frac{x_i - \mathbb{E}[x_i]}{\sqrt{\text{var}(x_i)}} + \beta_i$$

These allow to represent the identity $x_i \mapsto x_i$, if that was the optimal thing to do in the original network.

Ioffe and Szegedy, "Batch normalization: Accelerating deep network training by reducing internal covariate shift", ICML 2015

Batch normalization: Using mini-batches

Avoid analyzing the entire training set at each parameter update.

Input: Values of x over a mini-batch: $\mathcal{B} = \{x_{1\dots m}\}$;

Parameters to be learned: γ, β

Output: $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{scale and shift}$$

Ioffe and Szegedy, "Batch normalization: Accelerating deep network training by reducing internal covariate shift", ICML 2015

Batch normalization: Using mini-batches

Avoid analyzing the entire training set at each parameter update.

Input: Values of x over a mini-batch: $\mathcal{B} = \{x_{1\dots m}\}$;

Parameters to be learned: γ, β

Output: $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{scale and shift}$$

The batchnorm transformation makes each training example interact with the **other examples** in each mini-batch.

Ioffe and Szegedy, "Batch normalization: Accelerating deep network training by reducing internal covariate shift", ICML 2015

Batch normalization: Properties

Typically, batchnorm is applied right before the nonlinearity:

$$\sigma(\mathbf{W}\mathbf{x} + \mathbf{b}) \text{ becomes } \sigma \circ \text{BN}_{\gamma, \beta}(\mathbf{W}\mathbf{x})$$

The **bias** can be removed, since it is ruled out by the mean subtraction.

Ioffe and Szegedy, "Batch normalization: Accelerating deep network training by reducing internal covariate shift", ICML 2015

Batch normalization: Properties

Typically, batchnorm is applied right before the nonlinearity:

$$\sigma(\mathbf{W}\mathbf{x} + \mathbf{b}) \text{ becomes } \sigma \circ \text{BN}_{\gamma, \beta}(\mathbf{W}\mathbf{x})$$

The **bias** can be removed, since it is ruled out by the mean subtraction.

At **test** time, mini-batches are not legitimate.

Ioffe and Szegedy, "Batch normalization: Accelerating deep network training by reducing internal covariate shift", ICML 2015

Batch normalization: Properties

Typically, batchnorm is applied right before the nonlinearity:

$$\sigma(\mathbf{W}\mathbf{x} + \mathbf{b}) \text{ becomes } \sigma \circ \text{BN}_{\gamma, \beta}(\mathbf{W}\mathbf{x})$$

The **bias** can be removed, since it is ruled out by the mean subtraction.

At **test** time, mini-batches are not legitimate. Mean and variance are those estimated during training, and used for inference.

Ioffe and Szegedy, "Batch normalization: Accelerating deep network training by reducing internal covariate shift", ICML 2015

Batch normalization: Properties

Typically, batchnorm is applied right before the nonlinearity:

$$\sigma(\mathbf{W}\mathbf{x} + \mathbf{b}) \text{ becomes } \sigma \circ \text{BN}_{\gamma, \beta}(\mathbf{W}\mathbf{x})$$

The **bias** can be removed, since it is ruled out by the mean subtraction.

At **test** time, mini-batches are not legitimate. Mean and variance are those estimated during training, and used for inference.

Benefits:

- The stochastic uncertainty of the batch statistics acts as a **regularizer** that can benefit generalization.

Ioffe and Szegedy, "Batch normalization: Accelerating deep network training by reducing internal covariate shift", ICML 2015; Santurkar et al, "How does batch normalization help optimization?", NIPS 2018

Batch normalization: Properties

Typically, batchnorm is applied right before the nonlinearity:

$$\sigma(\mathbf{W}\mathbf{x} + \mathbf{b}) \text{ becomes } \sigma \circ \text{BN}_{\gamma, \beta}(\mathbf{W}\mathbf{x})$$

The **bias** can be removed, since it is ruled out by the mean subtraction.

At **test** time, mini-batches are not legitimate. Mean and variance are those estimated during training, and used for inference.

Benefits:

- The stochastic uncertainty of the batch statistics acts as a **regularizer** that can benefit generalization.
- Batchnorm leads to more **stable gradients**, thus **faster training** can be achieved with higher learning rates.

Ioffe and Szegedy, "Batch normalization: Accelerating deep network training by reducing internal covariate shift", ICML 2015; Santurkar et al, "How does batch normalization help optimization?", NIPS 2018

Normalization variants

Normalizing along the **batch dimension** can lead to inconsistency:

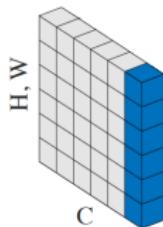
- Bad transfer across different data distributions.
- Reducing the **mini-batch size** increases error.

Normalization variants

Normalizing along the **batch dimension** can lead to inconsistency:

- Bad transfer across different data distributions.
- Reducing the **mini-batch size** increases error.

Several variants:

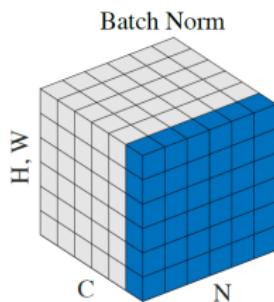


Normalization variants

Normalizing along the **batch dimension** can lead to inconsistency:

- Bad transfer across different data distributions.
- Reducing the **mini-batch size** increases error.

Several variants:

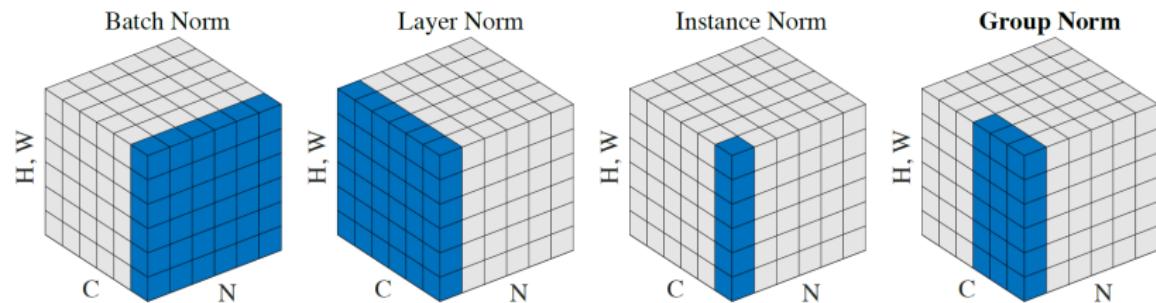


Normalization variants

Normalizing along the **batch dimension** can lead to inconsistency:

- Bad transfer across different data distributions.
- Reducing the **mini-batch size** increases error.

Several variants:



Ensemble deep learning?

Assume you have unlimited computational power.

Train an **ensemble** of deep nets and average their predictions.

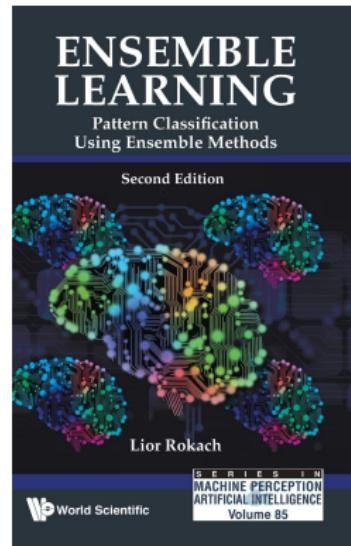
Ensemble deep learning?

Assume you have unlimited computational power.

Train an **ensemble** of deep nets and average their predictions.

Ensemble predictions (e.g. bayesian networks, random forests) are known to generalize better than the individual models.

Most successful methods in **Kaggle** are ensemble methods.



Ensemble deep learning?

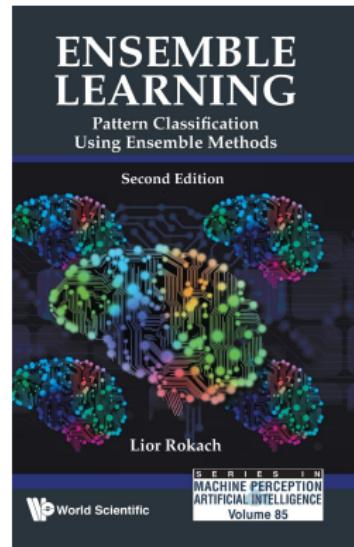
Assume you have unlimited computational power.

Train an **ensemble** of deep nets and average their predictions.

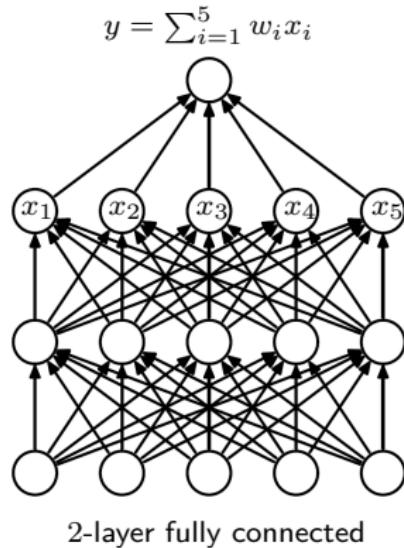
Ensemble predictions (e.g. bayesian networks, random forests) are known to generalize better than the individual models.

Most successful methods in **Kaggle** are ensemble methods.

However, for deep nets this would come at a **high computational cost**.

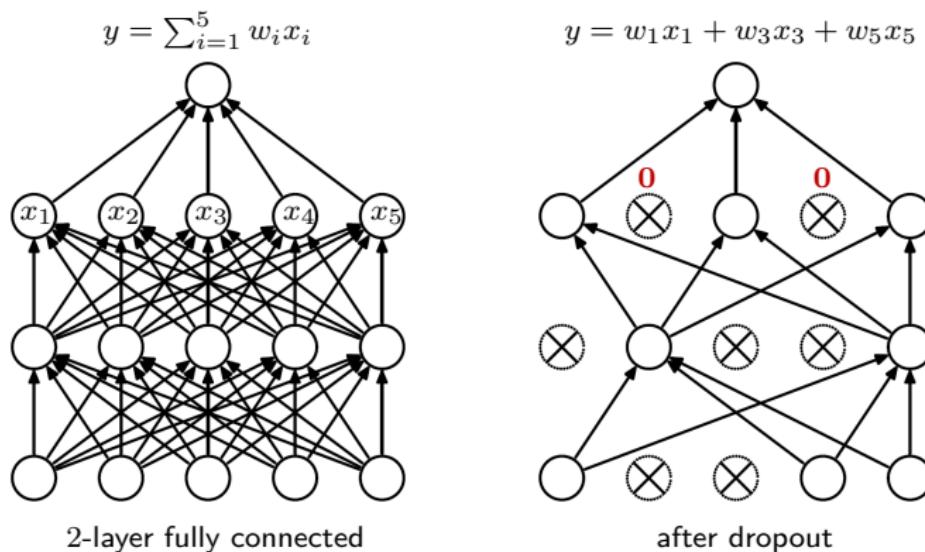


Dropout



Dropout

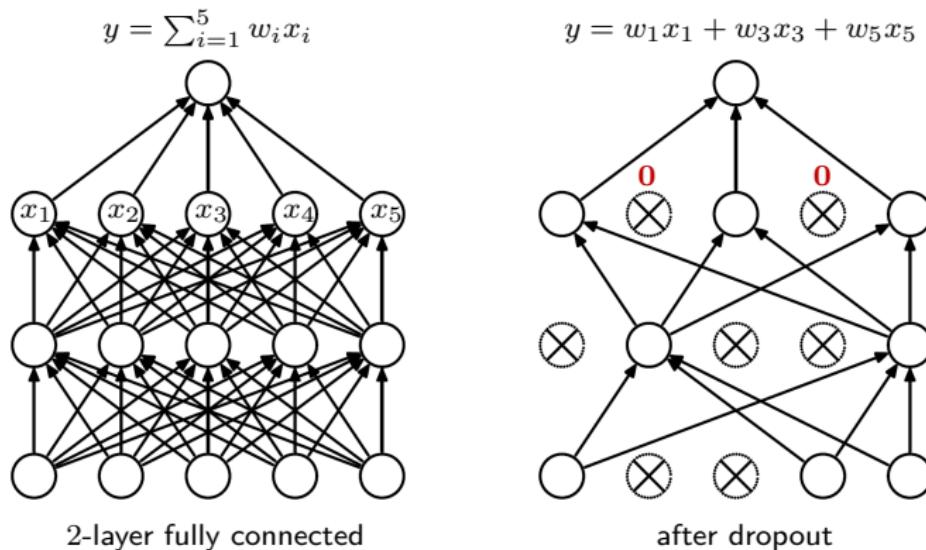
Main idea: Parametrize each model in the ensemble by dropping random units (i.e. nodes with their input/output connections):



Srivastava et al, "Dropout: A Simple Way to Prevent Neural Networks from Overfitting", JMLR 2014

Dropout

Main idea: Parametrize each model in the ensemble by **dropping** random units (i.e. nodes with their input/output connections):



Crucially, all networks **share** the same parameters.

Srivastava et al, “Dropout: A Simple Way to Prevent Neural Networks from Overfitting”, JMLR 2014

Dropout

Can be seen as **sampling** a network w.r.t. a probability distribution.

n nodes $\Rightarrow 2^n$ possible ways to sample them

Dropout

Can be seen as **sampling** a network w.r.t. a probability distribution.

n nodes $\Rightarrow 2^n$ possible ways to sample them

This is way too costly.

- **Training:** All the networks must be trained.
- **Test:** All the predictions must be averaged.

Srivastava et al, "Dropout: A Simple Way to Prevent Neural Networks from Overfitting", JMLR 2014

Dropout

Can be seen as **sampling** a network w.r.t. a probability distribution.

n nodes $\Rightarrow 2^n$ possible ways to sample them

Make it feasible by **keeping one single network**:

- **Training:** Generate a new sampling each time new training data is presented (e.g. **at each mini-batch** in SGD).

Dropout

Can be seen as **sampling** a network w.r.t. a probability distribution.

n nodes $\Rightarrow 2^n$ possible ways to sample them

Make it feasible by **keeping one single network**:

- **Training:** Generate a new sampling each time new training data is presented (e.g. **at each mini-batch** in SGD).

The **ensemble** is trained to convergence (e.g. with early stopping).

The individual models are **not** trained to convergence.

Dropout

Can be seen as **sampling** a network w.r.t. a probability distribution.

n nodes $\Rightarrow 2^n$ possible ways to sample them

Make it feasible by **keeping one single network**:

- **Test:** The trained weights from each model in the ensemble must be **averaged** somehow.

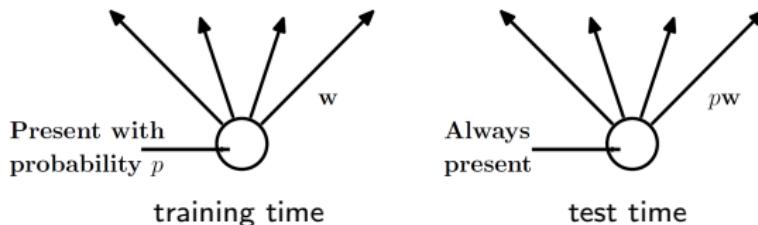
Dropout

Can be seen as **sampling** a network w.r.t. a probability distribution.

n nodes $\Rightarrow 2^n$ possible ways to sample them

Make it feasible by **keeping one single network**:

- **Test:** The trained weights from each model in the ensemble must be **averaged** somehow.



If a unit is retained with probability p during training (chosen by hand, even **per layer**), its outgoing weights are multiplied by p .

Srivastava et al, "Dropout: A Simple Way to Prevent Neural Networks from Overfitting", JMLR 2014

Dropout as an ensemble method

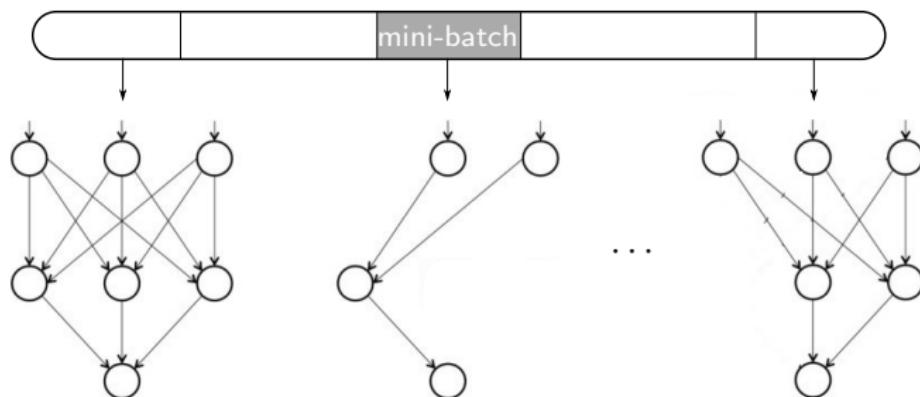
Dropout has two key features:

- It does **bagging**, i.e. each model is trained on random data.
- It does **weight sharing**, which is atypical in ensemble methods.

Dropout as an ensemble method

Dropout has two key features:

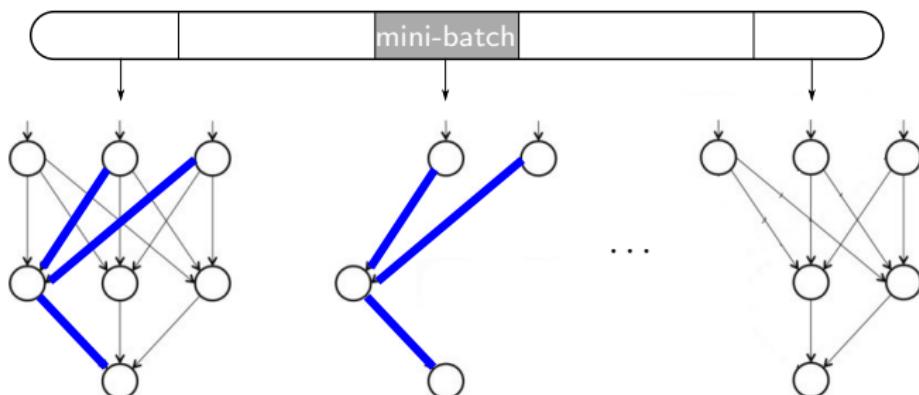
- It does **bagging**, i.e. each model is trained on random data.
- It does **weight sharing**, which is atypical in ensemble methods.



Dropout as an ensemble method

Dropout has two key features:

- It does **bagging**, i.e. each model is trained on random data.
- It does **weight sharing**, which is atypical in ensemble methods.



At each training step, the weight update is applied to all members of the ensemble simultaneously.

Dropout: Properties

In a standard neural network, weights are optimized **jointly**.

Co-adaptation: Small errors in a unit are absorbed by another unit.

Srivastava et al, "Dropout: A Simple Way to Prevent Neural Networks from Overfitting", JMLR 2014

Dropout: Properties

In a standard neural network, weights are optimized **jointly**.

Co-adaptation: Small errors in a unit are absorbed by another unit.

Some properties of dropout as a regularizer:

- Reduces co-adaptation by making units unreliable. This improves **generalization** to unseen data, and reduces overfitting.

Srivastava et al, "Dropout: A Simple Way to Prevent Neural Networks from Overfitting", JMLR 2014

Dropout: Properties

In a standard neural network, weights are optimized **jointly**.

Co-adaptation: Small errors in a unit are absorbed by another unit.

Some properties of dropout as a regularizer:

- Reduces co-adaptation by making units unreliable. This improves **generalization** to unseen data, and reduces overfitting.
- Side-effect: **sparse** representations are learned.

Srivastava et al, "Dropout: A Simple Way to Prevent Neural Networks from Overfitting", JMLR 2014

Dropout: Properties

In a standard neural network, weights are optimized **jointly**.

Co-adaptation: Small errors in a unit are absorbed by another unit.

Some properties of dropout as a regularizer:

- Reduces co-adaptation by making units unreliable. This improves **generalization** to unseen data, and reduces overfitting.
- Side-effect: **sparse** representations are learned.
- Performs closely to **exact** model averaging over all 2^n models.

Srivastava et al, "Dropout: A Simple Way to Prevent Neural Networks from Overfitting", JMLR 2014; Warde-Farley et al, "An empirical analysis of dropout in piecewise linear networks", 2014

Dropout: Properties

In a standard neural network, weights are optimized **jointly**.

Co-adaptation: Small errors in a unit are absorbed by another unit.

Some properties of dropout as a regularizer:

- Reduces co-adaptation by making units unreliable. This improves **generalization** to unseen data, and reduces overfitting.
- Side-effect: **sparse** representations are learned.
- Performs closely to **exact** model averaging over all 2^n models.
- ...and much better if no **weight sharing** is done in the exact model.

Srivastava et al, "Dropout: A Simple Way to Prevent Neural Networks from Overfitting", JMLR 2014; Warde-Farley et al, "An empirical analysis of dropout in piecewise linear networks", 2014

Dropout: Properties

In a standard neural network, weights are optimized **jointly**.

Co-adaptation: Small errors in a unit are absorbed by another unit.

Some properties of dropout as a regularizer:

- Reduces co-adaptation by making units unreliable. This improves **generalization** to unseen data, and reduces overfitting.
- Side-effect: **sparse** representations are learned.
- Performs closely to **exact** model averaging over all 2^n models.
- ...and much better if no **weight sharing** is done in the exact model.
- **Longer** training times, since parameter updates are now noisier.

Srivastava et al, "Dropout: A Simple Way to Prevent Neural Networks from Overfitting", JMLR 2014; Warde-Farley et al, "An empirical analysis of dropout in piecewise linear networks", 2014

Dropout: Properties

In a standard neural network, weights are optimized **jointly**.

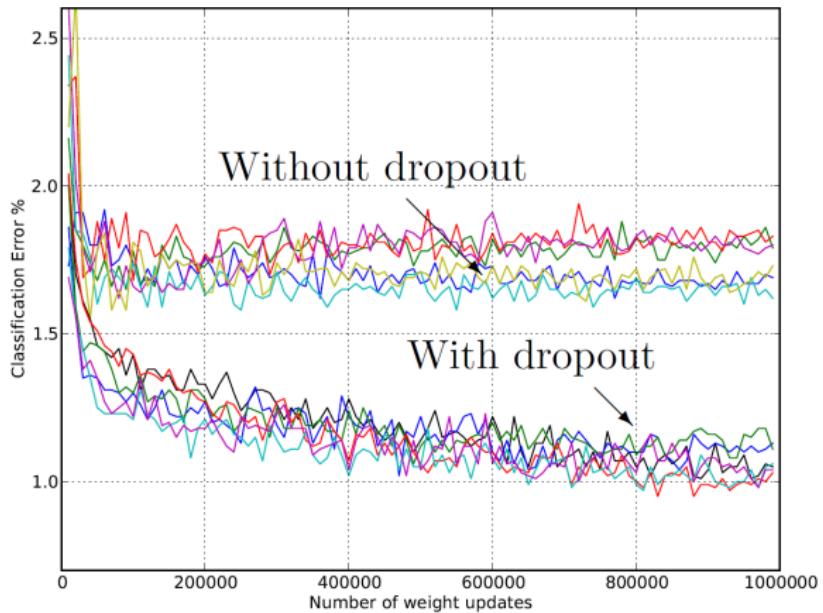
Co-adaptation: Small errors in a unit are absorbed by another unit.

Some properties of dropout as a regularizer:

- Reduces co-adaptation by making units unreliable. This improves **generalization** to unseen data, and reduces overfitting.
- Side-effect: **sparse** representations are learned.
- Performs closely to **exact** model averaging over all 2^n models.
- ...and much better if no **weight sharing** is done in the exact model.
- **Longer** training times, since parameter updates are now noisier.
- Typical choices: 20% of the input units and 50% of the hidden units.

Srivastava et al, "Dropout: A Simple Way to Prevent Neural Networks from Overfitting", JMLR 2014; Warde-Farley et al, "An empirical analysis of dropout in piecewise linear networks", 2014

Dropout: Properties



Srivastava et al, "Dropout: A Simple Way to Prevent Neural Networks from Overfitting", JMLR 2014

Suggested reading

- All the references given throughout the slides.
- Interesting thread on the history of double descent:
<https://twitter.com/hippopedoid/status/1243229021921579010>
- *Section 4.2.1* is a practical guide for batchnorm by the original authors:
<https://arxiv.org/pdf/1502.03167>
- *Appendix A* is a practical guide for dropout by the original authors:
<http://jmlr.org/papers/volume15/srivastava14a/srivastava14a.pdf>